Riccati eq., Linear Quadratic Regulator

Control Theory, Lecture 7

by Sergei Savin

Spring 2025

Content

- Hamilton-Jacobi-Bellman equation
  - Definitions
  - Cost, optimal cost
  - Differentiating optimal cost
- Algebraic Riccati equation
  - HJB for LTI
  - Linear Quadratic Regulator
  - Numerical methods

Let us define dynamics:

$$\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x}, \mathbf{u}) \tag{1}$$

with initial conditions $\mathbf{x}(0)$.

Additionally we define *control policy* as:

$$\mathbf{u} = \pi(\mathbf{x}, t) \tag{2}$$

To connect with the previous ways we talked about control, we can say that choosing different control gains and different feed-forward control amounts to choosing a different control policy.

Let J be an additive cost function:

$$J(\mathbf{x}_0, \pi(\mathbf{x}, t)) = \int_0^{\infty} g(\mathbf{x}, \mathbf{u}) dt \tag{3}$$

where $g(\mathbf{x}, \mathbf{u})$ is instantaneous cost and $\mathbf{x}_0 = \mathbf{x}(0)$ is the initial conditions. Notice that J depends on $\mathbf{x}_0$ rather than $\mathbf{x}(t)$, since initial conditions and control policy completely define the trajectory of the system $\mathbf{x}(t)$.

Let J^* be the optimal (lowest possible) cost. In other words:

$$J^*(\mathbf{x}_0) = \inf_{\pi} J(\mathbf{x}_0, \pi(\mathbf{x}, t)) \tag{4}$$

Optimal cost is attained when optimal policy is attained: $\pi = \pi^*(\mathbf{x}, t)$

With this, we can formulate *Hamilton-Jacobi-Bellman equation* (HJB):

$$\min_{\mathbf{u}} \left[g(\mathbf{x}, \mathbf{u}) + \frac{\partial J^*}{\partial \mathbf{x}} \mathbf{f}(\mathbf{x}, \mathbf{u}) \right] = 0 \tag{5}$$

We can find control that delivers minimum to the function (5):

$$\mathbf{u}^* = \underset{\mathbf{u}}{\operatorname{argmin}} \left[g(\mathbf{x}, \mathbf{u}) + \frac{\partial J^*}{\partial \mathbf{x}} \mathbf{f}(\mathbf{x}, \mathbf{u}) \right] \tag{6}$$

Algebraic Riccati (LTI case), 1

For LTI, dynamics is:

$$\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u} \tag{7}$$

We can choose quadratic cost:

$$g(\mathbf{x}, \mathbf{u}) = \mathbf{x}^\top \mathbf{Q}\mathbf{x} + \mathbf{u}^\top \mathbf{R}\mathbf{u} \tag{8}$$

Then HJB becomes:

$$\min_{\mathbf{u}} \left[\mathbf{x}^\top \mathbf{Q}\mathbf{x} + \mathbf{u}^\top \mathbf{R}\mathbf{u} + \frac{\partial J^*}{\partial \mathbf{x}}(\mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u})\right] = 0 \tag{9}$$

where $\mathbf{Q} = \mathbf{Q}^\top \geq 0$ and $\mathbf{R} = \mathbf{R}^\top > 0$.

Algebraic Riccati (LTI case), 2

There is a theorem that says that for LTI with quadratic cost, J^* has the form:

$$J^* = \mathbf{x}^\top \mathbf{S} \mathbf{x} \tag{10}$$

where $\mathbf{S} = \mathbf{S}^\top \geq 0$.

Then HJB becomes:

$$\min_{\mathbf{u}} \left[\mathbf{x}^\top \mathbf{Q} \mathbf{x} + \mathbf{u}^\top \mathbf{R} \mathbf{u} + \mathbf{x}^\top \mathbf{S}(\mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}) + (\mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u})^\top \mathbf{S} \mathbf{x} \right] = 0$$

Simplifying, we get:

$$\min_{\mathbf{u}} \left[\mathbf{u}^\top \mathbf{R} \mathbf{u} + \mathbf{x}^\top (\mathbf{Q} + \mathbf{S}\mathbf{A} + \mathbf{A}^\top \mathbf{S}) \mathbf{x} + \mathbf{x}^\top \mathbf{S} \mathbf{B} \mathbf{u} + \mathbf{u}^\top \mathbf{B}^\top \mathbf{S} \mathbf{x} \right] = 0$$

Finding partial derivative of the HJB with respect to $\mathbf{u}$ and setting it to zero (as it is an extreme point) we get:

$$2\mathbf{u}^\top \mathbf{R} + 2\mathbf{x}^\top \mathbf{S}\mathbf{B} = 0 \tag{11}$$

This expression can be transposed and $\mathbf{u}$ separated:

$$\mathbf{u} = -\mathbf{R}^{-1}\mathbf{B}^\top \mathbf{S}\mathbf{x} \tag{12}$$

This is the desired control law. We can see that it is *proportional*. We can re-write it as:

$$\mathbf{u} = -\mathbf{K}\mathbf{x} \tag{13}$$

where $\mathbf{K} = \mathbf{R}^{-1}\mathbf{B}^\top \mathbf{S}$ is the controller gain. This control law is called Linear Quadratic Regulator (LQR).

Algebraic Riccati (LTI case), 3

Substituting found control law into the HJB, we find:

$$\min_{\mathbf{u}} [\mathbf{x}^\top(\mathbf{Q} + \mathbf{SA} + \mathbf{A}^\top\mathbf{S})\mathbf{x} + \mathbf{x}^\top\mathbf{SBR}^{-1}\mathbf{RR}^{-1}\mathbf{B}^\top\mathbf{Sx} - \\ - \mathbf{x}^\top\mathbf{SBR}^{-1}\mathbf{B}^\top\mathbf{Sx} - \mathbf{x}^\top\mathbf{SBR}^{-1}\mathbf{B}^\top\mathbf{Sx}] = 0 \tag{14}$$

Simplifying, we get:

$$\mathbf{x}^\top(\mathbf{Q} + \mathbf{SA} + \mathbf{A}^\top\mathbf{S} - \mathbf{SBR}^{-1}\mathbf{B}^\top\mathbf{S})\mathbf{x} = 0 \tag{15}$$

which would hold for all $\mathbf{x}$ iff:

$$\mathbf{Q} - \mathbf{SBR}^{-1}\mathbf{B}^\top\mathbf{S} + \mathbf{SA} + \mathbf{A}^\top\mathbf{S} = 0 \tag{16}$$

This is the *Algebraic Riccati equation*.

LQR via software

There are a number of ways to solve LQR:

- In MATLAB there is a function `[K,S,P] = lqr(A,B,Q,R)`, where `P=eig(A-B*K)`
- In Python, there is `S = scipy.linalg.solve_continuous_are(A,B,Q,R)`

- Pole placement upsides: allows to design exactly how fast the control error decays to zero; allows to design control error oscillations.
- Pole placement downsides: may require unreasonably high control gains. Easy to ask for "unreasonable" performance.
- LQR upsides: easy to produce "reasonable" control gains.
- LQR downsides: may produce very slow decaying control error with oscillations.

Discrete case

Consider discrete dynamics:

$$\mathbf{x}_{i+1} = \mathbf{A}\mathbf{x}_i + \mathbf{B}\mathbf{u}_i \tag{17}$$

with a cost function:

$$J = \sum_{i=0}^{N} (\mathbf{x}_i^\top \mathbf{Q} \mathbf{x}_i + \mathbf{u}_i^\top \mathbf{R} \mathbf{u}_i) \tag{18}$$

Let us find the optimal control policy for this case.

Let us consider a linear control law $\mathbf{u}_i = -\mathbf{K}_i \mathbf{x}_i$. Given initial conditions $\mathbf{x}_0$ we find $\mathbf{x}_1$ as:

$$\mathbf{x}_1 = (\mathbf{A} - \mathbf{B}\mathbf{K}_i)\mathbf{x}_0 \tag{19}$$

let us define components of the cost associated with each iteration as J_1, J_2, etc:

$$J_i = \mathbf{x}_i^\top \mathbf{Q} \mathbf{x}_i + \mathbf{u}_i^\top \mathbf{R} \mathbf{u}_i \tag{20}$$

Since $\mathbf{u}_0 = -\mathbf{K}\mathbf{x}_0$ we can show that J_1 is a quadratic form of initial conditions:

$$J_0 = \mathbf{x}_0^\top (\mathbf{Q} + \mathbf{K}_i^\top \mathbf{R} \mathbf{K}_i)\mathbf{x}_0 \tag{21}$$

Similarly, we can show that all J_i can be written as a quadratic form of initial conditions.

Let us define cost-to-go as optimal cost for given initial conditions:

$$V_0 = \min_{\mathbf{u}} \sum_{i=0}^{N} (\mathbf{x}_i^\top \mathbf{Q} \mathbf{x}_i + \mathbf{u}_i^\top \mathbf{R} \mathbf{u}_i) \tag{22}$$

If $\mathbf{x}_0$, $\mathbf{x}_1$, $\mathbf{x}_2$, ... is a sequence of states that form an optimal trajectory, let us define the cost-to-go starting from each of these states as:

$$V_i = \min_{\mathbf{u}} \sum_{k=i}^{N} (\mathbf{x}_k^\top \mathbf{Q} \mathbf{x}_k + \mathbf{u}_k^\top \mathbf{R} \mathbf{u}_k) \tag{23}$$

We can note that the optimal cost will take a form of a quadratic function:

$$V_i = \mathbf{x}_i^\top \mathbf{P}_i \mathbf{x}_i \tag{24}$$

Cost-to-go, 2

Defining $\mathbf{x}_0 = \mathbf{z}$ and $\mathbf{u}_0 = \mathbf{w}$ we can write cost-to-go as:

$$V_0(\mathbf{z}) = \min_{\mathbf{w}} (\mathbf{z}^\top \mathbf{Q} \mathbf{z} + \mathbf{w}^\top \mathbf{R} \mathbf{w} + V_1(\mathbf{A}\mathbf{z} + \mathbf{B}\mathbf{w})) \qquad (25)$$

where $V_1(\mathbf{A}\mathbf{z} + \mathbf{B}\mathbf{w})$ is the optimal cost-to-go on the next step.

As the next step is closer to the goal, the optimal cost-to-go on the next step is both smaller than on the current step, and is contained in it.

The equation (25) is called *Bellman* equation.

Discrete LQR, 1

Since $V_0 = \mathbf{x}_0^\top \mathbf{P}_0 \mathbf{x}_0$ and $V_1 = \mathbf{x}_1^\top \mathbf{P}_1 \mathbf{x}_1$ we can re-write Bellman equation as:

$$\mathbf{z}^\top \mathbf{P}_0 \mathbf{z} = \min_{\mathbf{w}} (\mathbf{z}^\top \mathbf{Q} \mathbf{z} + \mathbf{w}^\top \mathbf{R} \mathbf{w} + (\mathbf{A}\mathbf{z} + \mathbf{B}\mathbf{w})^\top \mathbf{P}_1 (\mathbf{A}\mathbf{z} + \mathbf{B}\mathbf{w})) \quad (26)$$

To find minimum over $\mathbf{w}$ we take partial derivative and set it to zero:

$$\frac{\partial}{\partial \mathbf{w}} \Big(\mathbf{z}^\top \mathbf{Q} \mathbf{z} + \mathbf{w}^\top \mathbf{R} \mathbf{w} + (\mathbf{A}\mathbf{z} + \mathbf{B}\mathbf{w})^\top \mathbf{P}_1 (\mathbf{A}\mathbf{z} + \mathbf{B}\mathbf{w}) \Big) = 0 \quad (27)$$

$$2\mathbf{w}^\top \mathbf{R} + 2(\mathbf{A}\mathbf{z} + \mathbf{B}\mathbf{w})^\top \mathbf{P}_1 \mathbf{B} = 0 \quad (28)$$

$$\mathbf{R}\mathbf{w} + \mathbf{B}^\top \mathbf{P}_1 \mathbf{B}\mathbf{w} + \mathbf{B}^\top \mathbf{P}_1 \mathbf{A}\mathbf{z} = 0 \quad (29)$$

From $\mathbf{R}\mathbf{w} + \mathbf{B}^\top \mathbf{P}_1 \mathbf{B}\mathbf{w} + \mathbf{B}^\top \mathbf{P}_1 \mathbf{A}\mathbf{z} = 0$, we can find expression for the optimal control law:

$$(\mathbf{R} + \mathbf{B}^\top \mathbf{P}_1 \mathbf{B})\mathbf{w} = -\mathbf{B}^\top \mathbf{P}_1 \mathbf{A}\mathbf{z} \tag{30}$$

$$\mathbf{w} = -(\mathbf{R} + \mathbf{B}^\top \mathbf{P}_1 \mathbf{B})^{-1}\mathbf{B}^\top \mathbf{P}_1 \mathbf{A}\mathbf{z} \tag{31}$$

In general, given optimal cost-to-go matrix $\mathbf{P}_{i+1}$we find optimal control law:

$$\mathbf{u}_i = -(\mathbf{R} + \mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{B})^{-1}\mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{A}\mathbf{x}_i \tag{32}$$

Let us define $\mathbf{M} = (\mathbf{R} + \mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{B})^{-1}$ and $\mathbf{N} = \mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{A}$ we can re-write the control law:

$$\mathbf{u}_i = -\mathbf{M}\mathbf{N}\mathbf{x}_i \tag{33}$$

We can substitute the control law into the Bellman eq.:

$$\mathbf{x}_i^\top \mathbf{P}_i \mathbf{x}_i = \min_{\mathbf{u}} (\mathbf{x}_i^\top \mathbf{Q}\mathbf{x}_i + \mathbf{u}_i^\top \mathbf{R}\mathbf{u}_i + (\mathbf{A}\mathbf{x}_i + \mathbf{B}\mathbf{u}_i)^\top \mathbf{P}_{i+1} (\mathbf{A}\mathbf{x}_i + \mathbf{B}\mathbf{u}_i))$$

$$\begin{aligned} \mathbf{x}_i^\top \mathbf{P}_i \mathbf{x}_i = \mathbf{x}_i^\top \mathbf{Q}\mathbf{x}_i + \mathbf{x}_i^\top \mathbf{N}^\top \mathbf{M}\mathbf{R}\mathbf{M}\mathbf{N}\mathbf{x}_i + \\ +(\mathbf{A}\mathbf{x}_i - \mathbf{B}\mathbf{M}\mathbf{N}\mathbf{x}_i)^\top \mathbf{P}_{i+1} (\mathbf{A}\mathbf{x}_i - \mathbf{B}\mathbf{M}\mathbf{N}\mathbf{x}_i) \end{aligned}$$

$$\begin{aligned} \mathbf{P}_i = \mathbf{Q} + \mathbf{N}^\top \mathbf{M}\mathbf{R}\mathbf{M}\mathbf{N} + \mathbf{A}^\top \mathbf{P}_{i+1} \mathbf{A} - \mathbf{A}^\top \mathbf{P}_{i+1} \mathbf{B}\mathbf{M}\mathbf{N} - \\ \mathbf{N}^\top \mathbf{M}\mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{A} + \mathbf{N}^\top \mathbf{M}\mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{B}\mathbf{M}\mathbf{N} \end{aligned}$$

$$\mathbf{P}_i = \mathbf{Q} + \mathbf{A}^\top \mathbf{P}_{i+1} \mathbf{A} + \mathbf{N}^\top \mathbf{M}(\mathbf{R} + \mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{B})\mathbf{M}\mathbf{N} - \mathbf{N}^\top \mathbf{M}\mathbf{N} - \mathbf{N}^\top \mathbf{M}\mathbf{N}$$

$$\mathbf{P}_i = \mathbf{Q} + \mathbf{A}^\top \mathbf{P}_{i+1} \mathbf{A} - \mathbf{N}^\top \mathbf{M}\mathbf{N}$$

From $\mathbf{P}_i = \mathbf{Q} + \mathbf{A}^\top \mathbf{P}_{i+1} \mathbf{A} - \mathbf{N}^\top \mathbf{M} \mathbf{N}$ we obtain the final result:

$$\mathbf{P}_i = \mathbf{Q} + \mathbf{A}^\top \mathbf{P}_{i+1} \mathbf{A} - \mathbf{A}^\top \mathbf{P}_{i+1} \mathbf{B} (\mathbf{R} + \mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{B})^{-1} \mathbf{B}^\top \mathbf{P}_{i+1} \mathbf{A}$$

This equation can be used to compute $\mathbf{P}_i$ from known $\mathbf{P}_{i+1}$.

Further reading

- Underactuated robotics. Linear Quadratic Regulators.
- Discrete LQR. Stanford, EE363.
- Discrete LQR (infinite horizon). Stanford, EE363.

Lecture slides are available via Github:

github.com/SergeiSa/Control-Theory-2025

Appendix A: Illustration of HJB

Consider the additive cost $J(\mathbf{x}_0, \pi(\mathbf{x})) = \int_0^\infty g(\mathbf{x}, \mathbf{u}) dt$, where $\mathbf{u} = \pi(\mathbf{x})$ is a control policy. The function $g(\mathbf{x}, \mathbf{u}) \geq 0$ can be interpreted as a rate of change of cost.

Let $\pi^*(\mathbf{x})$ be the optimal control policy. Applying the optimal policy to the dynamics $\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u})$ we obtain optimal dynamics:

$$\dot{\mathbf{x}} = f^*(\mathbf{x}) = f(\mathbf{x}, \pi^*(\mathbf{x})) \tag{34}$$

Given initial conditions $\mathbf{x}_0 = \mathbf{z}$ we generate an optimal trajectory $\mathbf{x}^* = \mathbf{x}^*(t, \mathbf{z})$. Given optimal trajectory and optimal control policy we find optimal cost:

$$J^*(\mathbf{z}) = J(\mathbf{z}, \pi^*(\mathbf{x})) \tag{35}$$

Equivalently, we find optimal instantenious cost:

$$g^*(\mathbf{x}) = g(\mathbf{x}, \pi^*(\mathbf{x})) \tag{36}$$

Incurred cost

Since optimal cost depends on initial consitions only, we can find a function $J^* = J^*(\mathbf{z})$ defined over $\mathbb{R}^n$.

Lets us consider a trajectory $\mathbf{x}^* = \mathbf{x}^*(t, \mathbf{z})$ and sequential points on this trajectory $\mathbf{x}_0$, $\mathbf{x}_1$, $\mathbf{x}_2$, etc. associated with the time points t_0, t_1, t_2, etc. We can define incurred cost (incurred while moving from the initial state $\mathbf{z} = \mathbf{x}_0$ to the given point) for each of these points: S_0, S_1, S_2, etc.:

$$S_i = \int_0^{t_i} g^*(\mathbf{x}) dt \qquad (37)$$

Since $g^*(\mathbf{x}) \geq 0$, we observe that $S_0 \leq S_1 \leq S_2 \leq ....$ Moving along a trajectory $\mathbf{x}^*(t, \mathbf{z})$ we incur monotonically increasing cost. We can describe it as a time function $S(t)$. The rate of increace of this function is given by instantenious cost $g^*(\mathbf{x})$.

We know that the optimal cost from the point $\mathbf{z}$ is given as $J^*(\mathbf{z})$. For each sequential point on the trajectory we can define *cost-to-go* V_i as a difference between the optimal cost and the incurred cost:

$$V_i = J^*(\mathbf{z}) - S_i \tag{38}$$

For a given trajectory, we can describe cost-to-go as a time function $V(t) = J^*(\mathbf{z}) - S(t)$. Where as $S(t)$ is monotonically increasing, the function $V(t)$ is monotonically decreasing, with a rate of change given as $-g^*(\mathbf{x})$.

Note that the cost-to-go can be equivalently found as:

$$V(t) = J^*(\mathbf{x}^*(t)) \tag{39}$$

since the optimal cost we incur by starting from the point $\mathbf{x}_i$ is equivalent to the cost "have left to incur" when we reach $\mathbf{x}_i$ from $\mathbf{x}_0$.

With that, we can make an observation: for an optimal policy we see the rate of change of the cost-to-go function equal to $-g^*(\mathbf{x})$. But this rate of change can be found by taking a derivative of $J^*(\mathbf{x})$ with respect to the vector field $\dot{\mathbf{x}} = f^*(\mathbf{x})$:

$$-g^*(\mathbf{x}) = \frac{\partial J^*}{\partial \mathbf{x}} f^*(\mathbf{x}) \tag{40}$$

For sub-optimal control policies, the incurred cost will outpace the decrease of the cost-to-go:

$$g(\mathbf{x}, \mathbf{u}) + \frac{\partial J^*}{\partial \mathbf{x}} f(\mathbf{x}, \mathbf{u}) \geq 0 \tag{41}$$

Optimal policy recovers the sought equality:

$$\min_{\mathbf{u}} \left[g(\mathbf{x}, \mathbf{u}) + \frac{\partial J^*}{\partial \mathbf{x}} \mathbf{f}(\mathbf{x}, \mathbf{u}) \right] = 0 \tag{42}$$